

ExpoDesignAI — Architecture Audit vs. Target Vision

Audited: 2026-09-29 · Scope: all source files, UI, tests, data, docs

Executive Summary

Your codebase is a **solid, working monolith** (~2,800 lines of Python + 100KB of frontend) that already covers ~40% of your target architecture. The good news: the foundations (auth, RBAC, audit, document ingestion, vision AI, CAD extraction, RAG, IFC viewer) are real and working. The gap: the **middle tiers** of your target architecture (AI Orchestrator, Compliance Engine, Validation Engine, Report Engine, Knowledge Engine hybrid retrieval) **do not exist yet** — the system goes straight from "user asks a question" to "LLM answers from RAG chunks."

TARGET ARCHITECTURE	CURRENT STATE
ENGINEER UI	✓ Exists (monolithic HTML + viewer.js)
AI ORCHESTRATOR	✗ Missing — no planner/router/verifier
REASONING AI	⚠ Partial — single LLM, no orchestration
VISION AI	✓ Working (qwen2.5vl:32b via Ollama)
CODING AI	✗ Missing — no code generation agent
DOCUMENT INTELLIGENCE	✓ Working (PDF/Excel/Word/CAD/Image)
KNOWLEDGE ENGINE	⚠ Partial — FAISS only, no hybrid retrieval
COMPLIANCE ENGINE	⚠ Scaffolded — checklist exists, no engine
VALIDATION ENGINE	✗ Missing — no PASS/FAIL/REVIEW pipeline
REPORT ENGINE	✗ Missing — deps installed, not wired
Drawing/IFC/Revit Engines	⚠ Partial — viewer-only, no Revit API

Layer-by-Layer Gap Analysis

1. ENGINEER UI — Chat + Files + Viewer

```
flowchart LR
    subgraph Current["Current State"]
        A["index.html<br/>57KB monolith"]
        B["viewer.html + viewer.js<br/>IFC viewer"]
        C["classic.html<br/>unused legacy"]
    end
    subgraph Target["Target Vision"]
        D["Unified workspace shell"]
        E["3-region layout"]
        F["Embedded viewer panel"]
    end
    A -.->|"Phase 1 done"| D
    B -.->|"Re-host in shell"| F
```

Feature	Status	Detail
Chat interface	✓ Done	SSE streaming, multi-file attach, paste screenshots
File upload + management	✓ Done	Categorised storage, progress tracking, preview
IFC 3D viewer	✓ Done	web-ifc + three.js, sections, measure, Ask AI
Unified workspace shell	△□ Started	Phase 1 was planned, some UI work done (see .bak files)
Drawing viewer panel	✗ Missing	PDFs open inline, no dedicated drawing review panel
Schedule viewer	✗ Missing	Excel data extracted but no dedicated schedule view
Report viewer	✗ Missing	No report display/generation UI

[!TIP] You have 30+ **.bak files** in `ui/` — evidence of heavy iteration. Consider putting `ui/` under version control (even a simple local git) instead of accumulating backup files.

2. AI ORCHESTRATOR — Plan • Route • Verify

Component	Status	Detail
Planner	✗ Missing	No task decomposition. User query goes directly to LLM
Router	✗ Missing	Single model (<code>qwen2.5v1:32b</code>) for all tasks. No routing logic
Verifier	✗ Missing	No answer verification, no citation checking, no hallucination guard
Mode dispatch	△□ Basic	<code>mode</code> param distinguishes <code>chat</code> vs <code>cowork</code> — but this is a prompt switch, not orchestration

[!IMPORTANT] **This is the single biggest gap.** Your target architecture shows a central orchestrator that Plans → Routes to specialized AIs → Verifies output. Today, `main.py` line 566–778 (`build_messages` → `ask_stream`) is a **direct prompt** + **RAG** → **LLM** pipeline with zero orchestration.

What an orchestrator would do that doesn't exist today:

1. Classify the query type (compliance check? drawing comparison? schedule review? general Q&A?)
2. Select which AI(s) to invoke (reasoning vs. vision vs. coding)
3. Assemble typed evidence from structured stores (not just text chunks)
4. Run deterministic calculations in code (area checks, counts, comparisons)
5. Verify the answer against source evidence before returning it

3. REASONING AI — GLM / Qwen

Component	Status	Detail
-----------	--------	--------

LLM integration	✓ Working	qwen2.5v1:32b via Ollama, streaming, retry logic
Model abstraction	✓ Good	local_chat.py is clean, model-agnostic, retry-aware
Multi-model support	✗ Missing	Hardcoded single model. No GLM, no Qwen routing, no model selection
Reasoning chains	✗ Missing	No chain-of-thought, no multi-step reasoning, no tool-calling protocol

[!NOTE] local_chat.py and local_embed.py are well-engineered abstractions. They're the natural place to add model routing — the socket is there, you just need the switchboard.

4. VISION AI — Drawing/PDF

Component	Status	Detail
Image/drawing reading	✓ Working	vision.py — verbatim transcription, [illegible] marking, temperature=0
PDF page vision	✓ Working	extract.py — smart mode skips text-rich pages, vision on drawing sheets
CAD rendering + vision	✓ Working	cad.py — ezdxf render → vision model
Vision serialization	✓ Good	Thread lock prevents GPU OOM from concurrent vision calls
Drawing comparison	✗ Missing	Can read one drawing; cannot compare two revisions
Drawing annotation	✗ Missing	No markup/redline capability

5. CODING AI — C# / Python

Component	Status	Detail
Code generation	✗ Missing	No dedicated coding agent
Cowork mode	△□ Basic	<tool> tag protocol for sandboxed file ops (list/read/write/mkdir) — but this is prompt-based, not a real agent framework
C# / Revit API	✗ Missing	No C# code generation or Revit API integration
Python scripting	✗ Missing	No automated script generation for calculations

6. DOCUMENT INTELLIGENCE

```

flowchart TD
    subgraph Current
        PDF_C["PDF<br/>PyMuPDF + vision<br/>✓ Working"]
        XLS_C["Excel<br/>openpyxl<br/>✓ Working"]
    end

```

```

DOC_C["Word<br/>python-docx<br/>✓ Working"]
CAD_C["CAD<br/>ezdxf + ODA<br/>✓ Working"]
IMG_C["Image<br/>vision model<br/>✓ Working"]
end
subgraph Missing
  OCR_M["OCR<br/>✗ No Docling/<br/>Tesseract"]
  ARROW_M["Arrow<br/>✗ No columnar<br/>format"]
  STRUCT_M["Structured<br/>extraction<br/>✗ No typed<br/>entitie
end

```

Component	Status	Detail
PDF extraction	✓ Working	PyMuPDF text + rendered page → vision. Smart mode. Per-page JSON
Excel parsing	✓ Working	openpyxl, sheet/row level. Indexed into FAISS
Word parsing	✓ Working	python-docx, paragraphs + tables
CAD (DWG/DXF)	✓ Working	ezdxf data + matplotlib render → vision
Image reading	✓ Working	Base64 → Ollama vision
OCR / Docling	✗ Missing	No dedicated OCR. Vision model does OCR-like work but isn't specialized
Arrow / columnar	✗ Missing	No Apache Arrow integration for schedule data
Structured entity extraction	✗ Missing	Text is chunked as prose, not parsed into typed objects (doors, walls, specs)

[!TIP] Your diagram shows **Docling** for PDF. Docling would give you **layout-aware** extraction (tables, headers, figures as first-class objects) vs. the current flat text + vision approach. Worth evaluating as an upgrade to `extract.py`.

7. KNOWLEDGE ENGINE — Hybrid Retrieval

Component	Status	Detail
Vector search	✓ Working	FAISS with bge-m3 embeddings, per-project + shared KB
Keyword/BM25 search	✗ Missing	No BM25 or full-text search. Pure vector similarity
Hybrid retrieval	✗ Missing	No fusion of vector + keyword + structured queries
Reranker	✗ Missing	No cross-encoder or reranking stage
Source precedence	✗ Missing	No configurable priority (spec > drawing > schedule)
Entity-aware retrieval	✗ Missing	Cannot query "all doors on Level 2" — only text similarity

[!WARNING] **Hybrid retrieval is critical for engineering compliance.** Pure vector search misses exact code clause numbers, specific dimension values,

and schedule row IDs. A keyword/BM25 layer alongside vector search is table stakes for your domain.

8. Requirement DB + Vector DB

Component	Status	Detail
Requirement DB (structured)	⚠️ Scaffolded	findings table exists but is empty/unused. No qa_finding, bim_element, drawing, schedule_row
Vector DB	✓ Working (FAISS)	Per-project FAISS indexes. In-memory cache. Disk persistence
Persistent vector DB	✗ Not migrated	Chroma is in requirements.txt but completely unused . FAISS is the live system
Project-partitioned vectors	✓ Done	Each project has its own FAISS index
Knowledge graph	✗ Missing	No graph_edge table. No entity-relation navigation

9. COMPLIANCE ENGINE — Rules + AI

Component	Status	Detail
Rule definitions	⚠️ Scaffolded	disciplines.py has 25 Architecture checklist items with IBC/DBC code references
Deterministic rule checking	✗ Missing	No code compares actual vs. required values. LLM does all "checking"
AI-augmented compliance	⚠️ Prompt-only	Reviewer prompt tells LLM to check compliance, but there's no structured pipeline
Multi-discipline	⚠️ Scaffolded	Structural, MEP, Landscape, BIM, Cost/VE are registered but have "status": "planned", empty checklists
Code knowledge base	✓ Working	__codes__ FAISS collection with DBC/IBC excerpts

[!IMPORTANT] The gap between your target and current state is **deterministic vs. probabilistic compliance checking**. Your target shows a Compliance Engine with "Rules + AI" — implying rule-based checks (code in Python) augmented by AI interpretation. Today, the LLM does 100% of the compliance reasoning, which is unreliable for safety-critical engineering reviews.

10. Drawing Engine / IFC Engine / Revit API

Component	Status	Detail
Drawing Engine	⚠️ Read-only	Can extract and read drawings. Cannot generate, annotate, or compare
IFC Engine (browser)	✓ Working	web-ifc + three.js, full viewer with sections, measure, properties

IFC Engine (server)	✗ Missing	No IfcOpenShell. No server-side BIM entity extraction
Revit API	✗ Missing	No Revit integration. No .rvt parsing. No MCP server connection
BIM data storage	✗ Missing	No <code>bim_model</code> or <code>bim_element</code> tables

11. VALIDATION ENGINE — PASS / FAIL / REVIEW

Component	Status	Detail
Structured validation pipeline	✗ Missing	No automated PASS/FAIL/REVIEW workflow
Finding creation	⚠ Schema only	<code>findings</code> table exists (id, project, code, status, severity, comment) but nothing writes to it
Finding status tracking	✗ Missing	No lifecycle (open → pass/fail → resolved)
Automated QA runs	✗ Missing	No <code>/qa/run</code> endpoint, no batch checking

12. REPORT ENGINE — PDF / Excel / 3D

Component	Status	Detail
PDF reports	✗ Not built	<code>reportlab</code> is installed but unused
Excel reports	✗ Not built	<code>openpyxl</code> is installed (used for reading) but no report generation
Word reports	✗ Not built	<code>python-docx</code> is installed but no report generation
HTML templates	✗ Not built	<code>jinja2</code> is installed but no report templates exist
3D annotated views	✗ Not built	Viewer is read-only, no export

Architecture Health Scorecard

Dimension	Score	Assessment
Code quality	● 7/10	Well-structured modules, good separation, retry logic, error handling. Docstrings present. <code>main.py</code> is too large (1,406 lines) but organized
Security	● 7/10	JWT auth, RBAC, folder isolation, path traversal protection, sandboxed cowork, audit trail. Default <code>admin/admin</code> password is a risk
Scalability	● 3/10	Single process, in-memory FAISS, SQLite single-writer. Hard ceiling at ~5 concurrent users
AI Architecture	● 4/10	Working LLM + vision + embeddings. But no orchestration, no routing, no verification, no deterministic checks

Data	●	SQLite with 9 tables covers auth/chat/docs/audit well. Missing
Architecture	5/10	structured BIM/drawing/schedule/QA tables
Document	●	Excellent deep extraction pipeline. Missing OCR specialization
Intelligence	7/10	and structured entity parsing
Compliance	●	Checklist scaffolded. No engine. LLM does all reasoning.
	2/10	Unreliable for safety-critical work
Reporting	●	Dependencies installed, zero implementation
	1/10	
Frontend	●	Functional but monolithic. 57KB single HTML file. No component
	5/10	architecture. 30+ backup files
DevOps	●	No CI/CD, no Docker, no version control visible, backup files
	2/10	instead of git, no automated testing pipeline

Priority Recommendations

● Critical — Build These First

1. AI Orchestrator (your central gap)

```
# Conceptual architecture for the orchestrator
class QueryOrchestrator:
    def process(self, query, project, user):
        # 1. PLAN - Classify the query
        intent = self.classify(query) # compliance_check | comparison

        # 2. ROUTE - Select evidence sources
        if intent == "compliance_check":
            evidence = self.gather_compliance_evidence(query, project)
            # Structured: code clause + project value + deterministic c
        elif intent == "comparison":
            evidence = self.gather_comparison_evidence(query, project)
            # Two revisions side-by-side

        # 3. REASON - Deterministic first, AI second
        result = self.deterministic_check(evidence) # code comparison
        explanation = self.llm_explain(result, evidence) # natural lan

        # 4. VERIFY - Check citations exist in evidence
        verified = self.verify_citations(explanation, evidence)

        return StructuredAnswer(result, explanation, evidence, verified)
```

2. Compliance Engine with deterministic checks

Move from "LLM guesses compliance" to "code checks values, LLM explains findings":

```
# Example: Travel distance check (ARC-EGR-03)
def check_travel_distance(project_value_m, code_limit_m, sprinklered):
```

```

"""Deterministic. Returns PASS/FAIL/REVIEW with exact numbers."""
if project_value_m is None:
    return Finding(status="REVIEW", issue="Travel distance not docu
if project_value_m <= code_limit_m:
    return Finding(status="PASS", actual=project_value_m, required=
return Finding(status="FAIL", actual=project_value_m, required=code
                difference=project_value_m - code_limit_m)

```

3. Hybrid Retrieval (vector + keyword)

Add BM25 alongside FAISS:

```

from rank_bm25 import BM25Okapi

class HybridRetriever:
    def retrieve(self, query, project, k=8):
        vector_results = self.faiss_search(query, k*2)    # semantic
        bm25_results = self.bm25_search(query, k*2)       # keyword exac
        fused = self.reciprocal_rank_fusion(vector_results, bm25_result
        reranked = self.cross_encoder_rerank(query, fused[:k*3])
        return reranked[:k]

```

🔴 Important — Build These Next

4. Validation Engine

- Wire the existing findings table to actual checks
- Build `/api/v1/qa/run` endpoint that runs checklist items against project data
- Return structured PASS ✓ / FAIL ✗ / REVIEW △ results

5. Report Engine

- Use jinja2 templates + reportlab or weasyprint for PDF output
- Structure: project info → findings summary → per-item detail → evidence → recommendations
- Excel summary export via openpyxl

6. Structured Data Layer

- Add the tables proposed in your `ARCHITECTURE_ASSESSMENT.md` §10
- `bim_element`, `drawing`, `schedule_row`, `qa_finding`, `graph_edge`
- Make Alembic authoritative for migrations

🟢 Good to Have — Build Later

7. Scale Infrastructure

- SQLite → PostgreSQL (only needed at 50+ users)
- FAISS → Chroma/Qdrant (persistent, multi-process safe)
- Single process → Gunicorn multi-worker

- Request queue (Redis + Celery) for AI inference

8. Frontend Architecture

- Break `index.html` into components (even vanilla web components)
- Implement the unified workspace shell (Phase 1 of ARCHITECTURE_ASSESSMENT.md)
- Add proper state management

9. Revit API Integration

- Connect to your existing Revit API Lab codebase
- Server-side BIM entity extraction
- MCP server for live Revit model queries

Module Map: Current → Target

flowchart TD

```

subgraph Current["CURRENT (what exists)"]
    main["main.py<br/>1406 lines<br/>API + RAG + orchestration"]
    db_mod["db.py<br/>690 lines<br/>9 tables"]
    auth_mod["auth.py<br/>154 lines<br/>JWT + RBAC"]
    extract_mod["extract.py<br/>113 lines<br/>PDF/Excel/Word"]
    vision_mod["vision.py<br/>97 lines<br/>Ollama vision"]
    cad_mod["cad.py<br/>186 lines<br/>DWG/DXF"]
    disc_mod["disciplines.py<br/>156 lines<br/>Architecture checkli"]
    local_chat_mod["local_chat.py<br/>139 lines<br/>Ollama streamin"]
    local_embed_mod["local_embed.py<br/>87 lines<br/>Ollama embeddi"]
    meta_mod["meta_parse.py<br/>145 lines<br/>Drawing metadata"]
end

```

```

subgraph Target["TARGET (what to build)"]
    orchestrator["🟡 orchestrator.py<br/>AI Orchestrator"]
    compliance["🟡 compliance.py<br/>Compliance Engine"]
    validation["🟡 validation.py<br/>Validation Engine"]
    reports["🟡 reports.py<br/>Report Engine"]
    hybrid["🟡 hybrid_retrieval.py<br/>Knowledge Engine"]
    model_router["🟡 model_router.py<br/>AI Gateway"]
    bim_service["🟡 bim_service.py<br/>BIM Data Layer"]
    qa_engine["🟡 qa_engine.py<br/>QA Engine"]
end

```

```

main -->|"extract routing logic"| orchestrator
main -->|"extract RAG logic"| hybrid
local_chat_mod -->|"wrap in"| model_router
local_embed_mod -->|"wrap in"| model_router
disc_mod -->|"powers"| compliance
db_mod -->|"extend with new tables"| bim_service
extract_mod -->|"feeds"| hybrid
vision_mod -->|"used by"| orchestrator

```

```
style orchestrator fill:#ff6b6b,stroke:#333,color:#fff
style compliance fill:#ff6b6b,stroke:#333,color:#fff
style hybrid fill:#ff6b6b,stroke:#333,color:#fff
style validation fill:#ffd93d,stroke:#333
style reports fill:#ffd93d,stroke:#333
style model_router fill:#ffd93d,stroke:#333
style bim_service fill:#ffd93d,stroke:#333
style qa_engine fill:#ffd93d,stroke:#333
```

Quick Wins (can do today, high impact)

1. **Add BM25 to retrieval** — `pip install rank-bm25`, 50 lines of code, immediate accuracy improvement for clause-number and dimension queries
 2. **Wire findings table** — The schema exists. Add one endpoint to POST findings from chat and one to GET them. Instant structured output
 3. **Break `main.py`** — Extract routing logic into `routers/` (FastAPI routers). Zero behavior change, massive maintainability gain
 4. **Git init** — Replace 30+ `.bak` files with proper version control. This is your most impactful 5-minute fix
 5. **Reranker** — Add a small cross-encoder model (e.g. `bge-reranker-v2-m3` in Ollama) to rerank RAG results. ~30 lines of code, measurable answer quality improvement
-

[!CAUTION] **Do not attempt a full rewrite.** Your existing code works, is well-understood, and has been battle-tested with real project data. The target architecture should be built **incrementally on top of** what exists, exactly as your `ARCHITECTURE_ASSESSMENT.md` recommends. Each new layer (orchestrator, compliance engine, etc.) should be a new Python module that the existing `main.py` calls into.